

- Share Drive folder to IBM group?
  - Can see revisions, cuts down on email chains?
  - Allows for inline commenting?
  - Yes, lets do this
    - Use Lance and Adam's gmail account
- Go over design document
  - Good start
  - Have more than one
    - Use the existing one to focus on infrastructure
    - Then make another one that follows the same template about Hadoop
      - Will have high-level components
      - Purposes of Hadoop
      - Main nodes, other requirements
      - Google'd getting started with Hadoop
        - Maybe have someone going through single node setup
        - Have someone else doing cluster setup
        - We're going to need 4-5 different OpenStack images to manage Hadoop, so that's why we want to evaluate what Hadoop will require
          - This will allow us to know what to put into the OS images
    - Hadoop installs using crowbar
      - Based on Chet
      - Used by Dell and made specifically to do these deployments
      - Bar-Clamp modules that plug into crowbar
      - Use this as a last resort since it's made by Dell
- Go over use cases
- Further discussion of content in Google Groups.
  - Linux Lab and what we want to implement
  - Hadoop
    - Seems to be the most interesting
    - OpenStack Heat
      - Orchestrated deployment of Hadoop
      - Michael's never used it
      - Ties in with Puppet and Chef
      - Allows us to build off of Amazon API's as templates
      - Ways to solve shared storage
        - Create different images
        - Use different storage types and have multiple machines create a virtual file system
    - This would be something very good for us to implement
      - Good experience for both us students and IBM
    - Figure out how fast to deploy this, we're going to be kinda on our own.

- This will be challenging as both sides don't have a lot of Hadoop knowledge
- What have we defined as tasks?
  - Send over HTML of project plan
  - Wants to review that our outline make sense?
  - Have clearly defined tasks
  - Have them prioritized
  - Try to be as complete as possible early on
- Try to get design document mostly completed for Hadoop in the next two evenings
- By Thursday night try to have tasks defined
  - Have CS dept auth working
  - Deploying Hadoop workload
    - Has sub-tasks, so lets define those
  - What are my user's trying to do? Find tasks from that.
  - Try to send this out at the end of the night Thursday so they can be reviewed